

# Optimization Models in Machine Learning: Introduction and Examples

University of Washington

July 30, 2020

# Outline

1. Intro to Optimization + Notebook 1
2. Linear Regression and Regularization + Notebook 2
3. Logistic Regression + Notebook 3
4. Outlier Removal + Notebook 4

## Optimization: Overview

A general optimization problem has the form

$$\begin{array}{ll}\underset{x}{\text{minimize}} & f_0(x) \\ \text{subject to} & f_i(x) \leq b_i, \quad i = 1, \dots, m,\end{array}$$

with components

- ▶  $x = (x_1, \dots, x_n)$  - optimization variable
- ▶  $f_0 : \mathbf{R}^n \rightarrow \mathbf{R}$  - objective function
- ▶  $f_i : \mathbf{R}^n \rightarrow \mathbf{R}$  - constraint functions;  $b_i$  - constraint bounds

# Optimization: Overview

A general optimization problem has the form

$$\begin{array}{ll}\text{minimize} & f_0(x) \\ \text{subject to} & f_i(x) \leq b_i, \quad i = 1, \dots, m,\end{array}$$

with components

- ▶  $x = (x_1, \dots, x_n)$  - optimization variable
- ▶  $f_0 : \mathbf{R}^n \rightarrow \mathbf{R}$  - objective function
- ▶  $f_i : \mathbf{R}^n \rightarrow \mathbf{R}$  - constraint functions;  $b_i$  - constraint bounds

Many applications:

- |                               |                              |
|-------------------------------|------------------------------|
| ▶ Data fitting and regression | ▶ Recommender systems        |
| ▶ Classification              | ▶ Optimal control            |
| ▶ Image processing            | ▶ Medical treatment planning |
| ▶ Portfolio optimization      |                              |

# Optimization: Problem Classes

There are different classes of optimization problems, which can determine a problem's difficulty and solution method:

- ▶ Constrained vs. Unconstrained
- ▶ Smooth vs. Nonsmooth
- ▶ Convex vs. Nonconvex

---

<sup>1</sup>Stephen Boyd and Lieven Vandenberghe. Convex Optimization. Cambridge University Press, 2004.

# Optimization: Problem Classes

There are different classes of optimization problems, which can determine a problem's difficulty and solution method:

- ▶ Constrained vs. Unconstrained
- ▶ Smooth vs. Nonsmooth
- ▶ Convex vs. Nonconvex

An important class: **convex optimization** problems

“With only a bit of exaggeration, we can say that if you formulate a practical problem as a convex optimization problem, then you have solved the original problem.”<sup>1</sup>

---

<sup>1</sup>Stephen Boyd and Lieven Vandenberghe. Convex Optimization. Cambridge University Press, 2004.

## Optimization: Problem Classes

A convex optimization problem has objective and constraint functions that satisfy the inequality

$$f_i(\lambda x + (1 - \lambda)y) \leq \lambda f_i(x) + (1 - \lambda)f_i(y)$$



for all  $x, y \in \mathbf{R}^n$  and all  $0 \leq \lambda \leq 1$ .

# Optimization: Problem Classes

A convex optimization problem has objective and constraint functions that satisfy the inequality

$$f_i(\lambda x + (1 - \lambda)y) \leq \lambda f_i(x) + (1 - \lambda)f_i(y)$$

for all  $x, y \in \mathbf{R}^n$  and all  $0 \leq \lambda \leq 1$ .



Important consequence: in a convex problem, no “local minima”



## Optimization: Solution Methods

Very few optimization problems have a closed-form solution (e.g., least-squares); most problems are solved using iterative methods.

One important iterative method is gradient descent:



$$x^{k+1} = x^k - \alpha \nabla f(x^k)$$

# Optimization: Solution Methods

Very few optimization problems have a closed-form solution (e.g., least-squares); most problems are solved using iterative methods.

One important iterative method is gradient descent:

$$x^{k+1} = x^k - \alpha \nabla f(x^k)$$



# Optimization: Machine Learning Models

Many problems in machine learning seek to build a model

$$g(a; x) \approx y$$

given a data set

$$\left\{ (a_1, y_1), \dots, (a_m, y_m) \right\},$$

with components

- ▶  $a_i = (a_{i1}, \dots, a_{in})$  - data features
- ▶  $y_i \in \mathbf{R}$  or  $\{0, 1\}$  - data value or label/class
- ▶  $g : \mathbf{R}^n \rightarrow \mathbf{R}$  or  $\{0, 1\}$  - prediction function
- ▶  $x = (x_1, \dots, x_n)$  - model parameters
- ▶  $m$  - number of data points
- ▶  $n$  - number of data features

# Optimization: Machine Learning Models

We can fit a model to the given data by solving an optimization problem of the form

$$\underset{x}{\text{minimize}} \quad \sum_{i=1}^m f_i(g(a_i; x), y_i) + r(x),$$


with components

- ▶  $x = (x_1, \dots, x_n)$  - model parameters we want to learn
- ▶  $f_i : \mathbf{R}^n \rightarrow \mathbf{R}$  - “loss” functions: measure how well the model fits the data for given parameters; e.g.,  $(g(a_i; x) - y_i)^2$
- ▶  $r(x) : \mathbf{R}^n \rightarrow \mathbf{R}$  - regularization function


# Optimization: Machine Learning Models

We focus on two common problems in machine learning:

## Linear Regression




## Logistic Regression (Classification)




# Optimization: Machine Learning Models

We focus on two common problems in machine learning:

## Linear Regression



## Logistic Regression (Classification)



- ▶ Data: Continuous features  $\{a_i\}$  and outputs  $\{y_i\}$
- ▶ Goal: Find linear predictor

$$x_0 + x_1 a_i \approx y_i$$

- ▶ Approach: Assume a statistical model for errors and develop a maximum likelihood formulation

## Linear Regression: Derivation

Assuming the errors in our data come from a normal distribution,

$$y_i = x_0 + x_1 a_i + \epsilon_i, \quad \epsilon_i \sim N(0, \sigma^2) \quad \text{independent,}$$

## Linear Regression: Derivation

Assuming the errors in our data come from a normal distribution,

$$y_i = x_0 + x_1 a_i + \epsilon_i, \quad \epsilon_i \sim N(0, \sigma^2) \quad \text{independent,}$$

the probability of observation  $(a_i, y_i)$  given the parameters is

$$P((a_i, y_i); x_0, x_1) \propto \exp\left(\frac{-(y_i - x_0 - x_1 a_i)^2}{2\sigma^2}\right).$$



# Linear Regression: Derivation

Assuming the errors in our data come from a normal distribution,



$$y_i = x_0 + x_1 a_i + \epsilon_i, \quad \epsilon_i \sim N(0, \sigma^2) \quad \text{independent,}$$

the probability of observation  $(a_i, y_i)$  given the parameters is

$$P((a_i, y_i); x_0, x_1) \propto \exp\left(\frac{-(y_i - x_0 - x_1 a_i)^2}{2\sigma^2}\right).$$


We can fit model parameters by maximizing the likelihood (minimizing the negative log-likelihood):

$$-\log \prod_{i=1}^m \exp\left(-\frac{(y_i - x_0 - x_1 a_i)^2}{2\sigma^2}\right) \propto \sum_{i=1}^m (y_i - x_0 - x_1 a_i)^2$$



# Linear Regression: Intuition and Properties

$$\min_{x_0, x_1} \sum_{i=1}^m (y_i - x_0 - x_1 a_i)^2$$



- ▶ Minimize the least-squares distance between observations  $y_i$  and predictions  $x_0 + x_1 a_i$ .
- ▶ The problem is convex, smooth, and easy to solve.
- ▶ Linear regression actually has a closed-form solution, but it is often found more efficiently by iterative algorithms

# Regularization: Overview

Many problems in machine learning add a regularization term  $r(x)$  to the objective function to

- ▶ incorporate prior knowledge about *structure* in  $x$ , e.g., sparsity or smoothness
- ▶ help avoid overfitting,
- ▶ get more robust (to data perturbations) solutions, or
- ▶ improve the stability of the solution process.

Two popular forms of regularized linear regression:



- ▶ Lasso -  $\min_x f(x) + \lambda \|x\|_1$ , where  $\|x\|_1 = \sum_{i=1}^n |x_i|$
- ▶ Ridge -  $\min_x f(x) + \lambda \|x\|_2^2$ , where  $\|x\|_2^2 = \sum_{i=1}^n x_i^2$

# Regularization: Geometric Interpretation

Consider the constrained least-squares problem

$$\begin{aligned} & \underset{x}{\text{minimize}} && \frac{1}{2} \|Ax - y\|_2^2 \\ & \text{subject to} && \|x\|_p \leq t. \end{aligned}$$

Choice of norm influences properties of solution  $x$ : with  $p = 1$ , solutions tend to occur on the vertices, where many  $x_i = 0$ .



## Regularization: Relaxed Constraints

We can move the norm from a constraint into the objective function to get

$$\underset{x}{\text{minimize}} \quad \frac{1}{2} \|Ax - y\|_2^2 + \lambda \|x\|_p,$$

where regularization parameter  $\lambda$  balances model error with how much we regularize.

The Lasso ( $p = 1$ ) is often used to find sparse solutions. Ridge regression ( $p = 2$ ) is often used for ill-conditioned problems.

More generally: regularizers can promote other structures: For example, if the parameters form a matrix  $X$ , a low-rank matrix is often desired (e.g., the ‘matrix completion problem’ for recommender systems).

# Logistic Regression: Overview

- ▶ Data: Continuous features  $\{a_i\}$  and discrete labels  $y_i \in \{0, 1\}$
- ▶ Goal: Find linear predictor

$$x_0 + x_1 a_i = \begin{cases} \text{positive} & \Rightarrow y_i = 1 \\ \text{negative} & \Rightarrow y_i = 0 \end{cases}$$

- ▶ Approach: Combine Bernoulli model with a linear predictor
- ▶ Examples: Hours studied vs. Pass/Fail, measurements vs. disease

## Logistic Regression: Derivation

Rewriting the Bernoulli model in standard form,

$$\begin{aligned}P\left((a_i, y_i); p_i\right) &= p_i^{y_i} (1 - p_i)^{1-y_i} \\&= \exp\left(y_i \log\left(\frac{p_i}{1 - p_i}\right) + \log(1 - p_i)\right),\end{aligned}$$

we can model the term multiplying  $y_i$  using our linear predictor,

$$\log\left(\frac{p_i}{1 - p_i}\right) = x_0 + x_1 a_i,$$

which gives us,

$$\log(1 - p_i) = -\log(1 + \exp(x_0 + x_1 a_i)).$$

Combining the above expressions results in the likelihood function

$$\mathcal{L}(x_0, x_1; (a, y)) = \prod_{i=1}^m \exp\left(y_i(x_0 + x_1 a_i) - \log(1 + \exp(x_0 + x_1 a_i))\right).$$

## Logistic Regression: Derivation

We can fit our model parameters to the given data by maximizing the likelihood, or by minimizing the negative log-likelihood:

$$-\log \mathcal{L}(x_0, x_1; (a, y)) = \sum_{i=1}^m \log(1 + \exp(x_0 + x_1 a_i)) - y_i(x_0 + x_1 a_i)$$

Explicitly, we solve the following problem


$$\min_{x_0, x_1} \sum_{i=1}^m \log(1 + \exp(x_0 + x_1 a_i)) - y_i(x_0 + x_1 a_i)$$



# Logistic Regression: Intuition and Properties


$$\min_{x_0, x_1} \sum_{i=1}^m \log(1 + \exp(x_0 + x_1 a_i)) - y_i(x_0 + x_1 a_i)$$

- ▶ If the label is 0, we want to make  $\log(1 + \exp(x_0 + x_1 a_i))$  as small as possible, equivalent to making  $x_0 + x_1 a_i \ll 0$
- ▶ If the label is 1, can show objective decreases with respect to  $x_0 + x_1 a_i$ , so we want  $x_0 + x_1 a_i \gg 0$



# Logistic Regression: Intuition and Properties

- We look for intercept  $x_0$  and slope  $x_1$  that do the best job for all the data in the set.



# Logistic Regression: Intuition and Properties

- ▶ The problem is convex and smooth, and 'nice' to solve.
- ▶ For a future data points with feature  $a$ ,  $p = \frac{\exp(x_0 + x_1 a)}{1 + \exp(x_0 + x_1 a)}$
- ▶ Other methods can also be used, e.g. support vector machines.

## Trimming: Outlier Removal

There are many classic ways to remove outliers:

# Trimming: Outlier Removal

There are many classic ways to remove outliers:

- ▶ Fit, remove outliers, refit
  - ▶ Upside: easy to do
  - ▶ Downside: outliers can affect the initial fit
  - ▶ Downside: when to stop?

# Trimming: Outlier Removal

There are many classic ways to remove outliers:

- ▶ Fit, remove outliers, refit
  - ▶ Upside: easy to do
  - ▶ Downside: outliers can affect the initial fit
  - ▶ Downside: when to stop?
- ▶ Regression: replace least squares loss with 'robust' penalty
  - ▶ Upside: relatively easy to do
  - ▶ Downside: how to pick shape?
  - ▶ Downside: how to extend to non-additive errors?

# Trimming: Outlier Removal

There are many classic ways to remove outliers:

- ▶ Fit, remove outliers, refit
  - ▶ Upside: easy to do
  - ▶ Downside: outliers can affect the initial fit
  - ▶ Downside: when to stop?
- ▶ Regression: replace least squares loss with 'robust' penalty
  - ▶ Upside: relatively easy to do
  - ▶ Downside: how to pick shape?
  - ▶ Downside: how to extend to non-additive errors?
- ▶ Our focus: trimming
  - ▶ Upside: works for any model
  - ▶ Upside: transparent assumptions
  - ▶ Downside: nonconvex model
  - ▶ Upside: doesn't seem to matter in practice

# Trimming: Overview

Trimming uses auxiliary weights to detect outliers:

$$\min_{x, w} \sum_{i=1}^m w_i f_i(x) \quad \text{s.t.} \quad w_i \in [0, 1], \quad \sum_{i=1}^m w_i = h$$

- ▶ For fixed  $x$ , minimal  $h$  residuals have  $w_i = 1$ , rest are 0
- ▶ Minimal  $h$  residuals are thus classified as 'inliers'
- ▶ Remaining  $m - h$  points are by default 'outliers'
- ▶ As  $x$  varies, we are looking to only fit inliers.

Problem is theoretically hard, but practically works very well.



## Trimming: Overview


The general idea extends to any learning model:

$$\min_{x, w} \sum_{i=1}^m w_i f_i(x) \quad \text{s.t.} \quad w_i \in [0, 1], \quad \sum_{i=1}^m w_i = h$$

- ▶ Least squares:  $f_i(x) = \frac{1}{2}(y_i - x_0 - x_1 a_i)^2$
- ▶ Logistic:  $f_i(x) = \log(1 + \exp(x_0 + x_1 a_i)) - y_i(x_0 + x_1 a_i)$
- ▶ Neural net:  $f_i(x) = \text{soft max}$  for a labeled data point

## Trimming: Least Squares Example

$$\min_{x, w} \sum_{i=1}^m \frac{w_i}{2} (y_i - x_0 - x_1 a_i)^2 \quad \text{s.t.} \quad w_i \in [0, 1], \quad \sum_{i=1}^m w_i = h$$



# Trimming: CNN Example

Here we see the results of a convolutional neural network (CNN) classifier that predicts cats and birds, with inliers (top row) and outliers (bottom row) identified using trimming.

